

浙江大学实验报告

课程名称: 计算机体系结构 实验类型: 综合
实验名称: Lab4: pipelined CPU with cache
指导教师: 何水兵 完成日期: 2023 年 11 月 14 日

实验分工

姓名: 蒋奕 3210103803

同组学生姓名: 黄俊涛 3210101831

分工情况:

蒋奕: cmu 模块 + 设计测试仿真

cmu 采用状态机实现, 最终通过仿真和上板子结果验证。

黄俊涛: cmu 模块 + 设计测试仿真

1 实验目的和要求

- Understand the principle of Cache Management Unit (CMU) and State Machine of CMU.
- Understand the principle of Cache Management Unit (CMU) and State Machine of CMU.
- Master verification methods of CMU and compare the performance of CPU when it has cache or not.
- Master the design methods of CMU and Integrate it to the CPU.

2 实验内容和原理

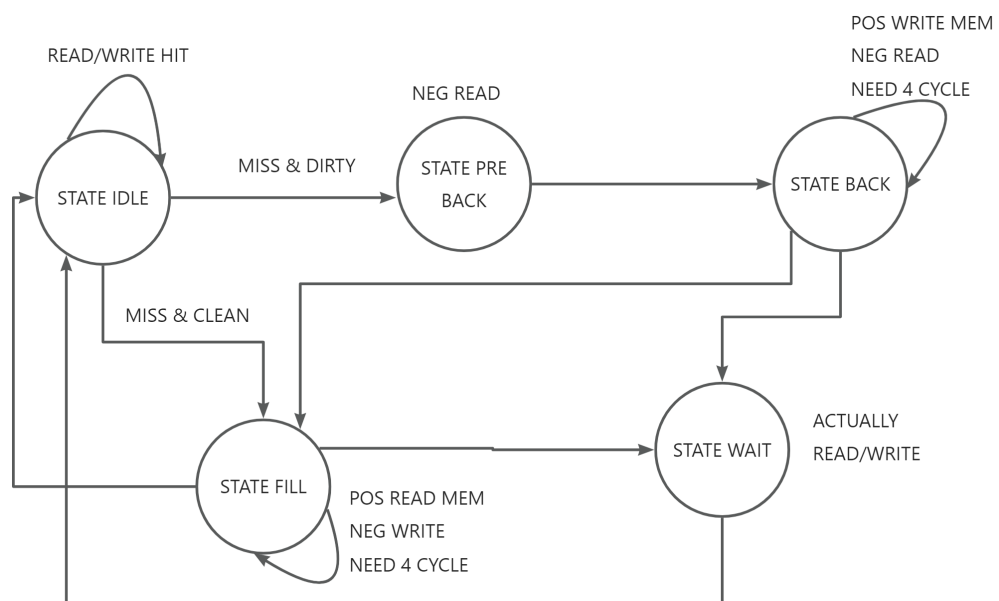
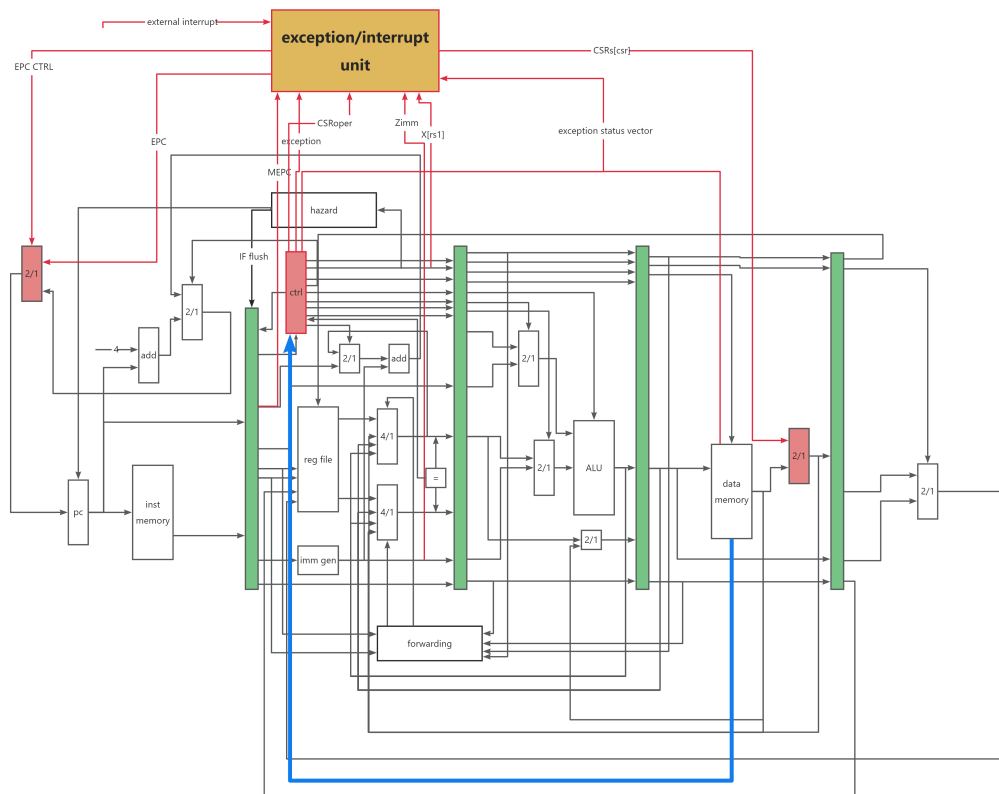
2.1 实验内容

- Design of Cache Management Unit and integrate it to CPU.
- Observe and Analyze the Waveform of Simulation.
- Compare the performance of CPU when it has cache or not.

2.2 实验原理

- Structure of Circuit

Compared to lab2, we add a line from data memory to ctrl unit. The circuit is as follows:



- Structure of State Machine

The schematic is as above:

Cache operations occur on the neg edge of the current state, while memory operations occur on the pos edge of the current state. Each state has meaning:

1. **S_IDLE:** Idle state; no memory operations. Cache operations continue if there is a hit, keeping the system in this state.
2. **S_PRE_BACK:** To facilitate writing back, perform a cache read operation.
3. **S_BACK:** On the rising edge, write back the data from the previous state to memory. On the falling edge, read the data to be written back from the cache. This process continues until the entire cache line is written back. Since memory operations take 4 cycles to complete, wait for the memory acknowledgment signal before changing the state.
4. **S_FILL:** On the rising edge, read data from memory. On the falling edge, write data to the cache. This process continues until the entire cache line is filled. Similar to S_BACK, wait for the memory acknowledgment signal.
5. **S_WAIT:** Execute cache operations that couldn't be performed earlier due to a miss.

3 实验过程和数据记录及结果分析

3.1 实验过程

The details of cache can be seen in comments of the corresponding code.

- state IDLE

```
S_IDLE: begin
    if (en_r || en_w) begin
        if (cache_hit)
            next_state = S_IDLE;
        else if (cache_valid && cache_dirty)
            next_state = S_PRE_BACK;
        else
            next_state = S_FILL;
    end
    next_word_count = 2'b00;
end
```

- state BACK

In this lab, in order to simulate Memory being slower than Cache, a counter has been manually set up. The state transition occurs only every 4 clock cycles.

```
S_BACK: begin
    if (mem_ack_i &&
```

```

word_count == {ELEMENT_WORDS_WIDTH{1'b1}})
    next_state = S_FILL;
else
    next_state = S_BACK;
if (mem_ack_i)
    next_word_count = word_count + 2'b01;
else
    next_word_count = word_count;
end

```

- state FILL

In this lab, in order to simulate Memory being slower than Cache, word count has been manually set up. The state transition occurs only every 4 clock cycles.

```

S_FILL: begin
    if (mem_ack_i &&
        word_count == {ELEMENT_WORDS_WIDTH{1'b1}})
        next_state = S_WAIT;
    else
        next_state = S_FILL;
    if (mem_ack_i)
        next_word_count = word_count + 2'b01;
    else
        next_word_count = word_count;
end

```

- state WAIT

```

S_WAIT: begin
    next_state = S_IDLE;
    next_word_count = 2'b00;
end

```

For the control part of the cache, we should do as follows:

- S_IDLE and S_WAIT

The cache operates normally. The cache_store signal is set to 1 because there is no need to transfer data from memory to the cache in these states.

```

S_IDLE, S_WAIT: begin
    cache_addr = addr_rw;
    cache_load = en_r;
    cache_edit = en_w;
    cache_store = 1'b0;
    cache_u_b_h_w = u_b_h_w;
end

```

```

        cache_din = data_w;

    end

```

- S_BACK, S_PRE_BACK

Due to the need to write back dirty data from the cache to memory, CPU should not operate on the cache during this time. Therefore, we ought to set `cache_load`, `cache_edit`, and `cache_store` to 0. Additionally, no data needs to be loaded into the cache during this operation.

```

S_BACK, S_PRE_BACK: begin
    cache_addr = {addr_rw[ADDR_BITS-1:BLOCK_WIDTH],
                  next_word_count,
                  {ELEMENT_WORDS_WIDTH{1'b0}}};
    cache_load = 1'b0;
    cache_edit = 1'b0;
    cache_store = 1'b0;
    cache_u_b_h_w = 3'b010;
    cache_din = 32'b0;

end

```

- S_FILL

In S_FILL, data is read from memory and placed into cache. To prevent unexpected issues like overwriting data, it's necessary to stop CPU from writing to cache or modifying data in cache. If `mem_ack_i` is 1, indicating that the cache has successfully read data from memory, the cache data can be written back to memory. Therefore, `cache_store` is set to 1. Since data is read from memory, `cache_din` is set to `mem_data_i`.

```

S_FILL: begin
    cache_addr = {addr_rw[ADDR_BITS-1:BLOCK_WIDTH],
                  word_count,
                  {ELEMENT_WORDS_WIDTH{1'b0}}};
    cache_load = 1'b0;
    cache_edit = 1'b0;
    cache_store = mem_ack_i;
    cache_u_b_h_w = 3'b010;
    cache_din = mem_data_i;

end

```

- stall signal

In CMU unit, stall signal is used to pause the execution of instructions in the CPU. When there is communication between the cache and memory, stall signal needs to be set to 1. Therefore, except for the S_IDLE state, a pause in the execution of instructions is required in all other states.

```

assign stall = (next_state != S_IDLE);

```

- Complete the cache.v file

```

1 module cmu (
2     // CPU side
3     input clk,
4     input rst,
5     input [31:0] addr_rw,
6     input en_r,
7     input en_w,
8     input [2:0] u_b_h_w,
9     input [31:0] data_w,
10    output [31:0] data_r,
11    output stall,
12
13    // mem side
14    output reg mem_cs_o = 0,
15    output reg mem_we_o = 0,
16    output reg [31:0] mem_addr_o = 0,
17    input [31:0] mem_data_i,
18    output [31:0] mem_data_o,
19    input mem_ack_i,
20
21    // debug info
22    output [2:0] cmu_state
23 );
24
25 `include "addr_define.vh"
26
27 reg [ADDR_BITS-1:0] cache_addr = 0;
28 reg cache_load = 0;
29 reg cache_store = 0;
30 reg cache_edit = 0;
31 reg [2:0] cache_u_b_h_w = 0;
32 reg [WORD_BITS-1:0] cache_din = 0;
33 wire cache_hit;
34 wire [WORD_BITS-1:0] cache_dout;
35 wire cache_valid;
36 wire cache_dirty;
37 wire [TAG_BITS-1:0] cache_tag;
38
39 cache CACHE (
40     .clk(~clk),

```

```

41         .rst(rst),
42         .addr(cache_addr),
43         .load(cache_load),
44         .store(cache_store),
45         .edit(cache_edit),
46         .invalid(1'b0),
47         .u_b_h_w(cache_u_b_h_w),
48         .din(cache_din),
49         .hit(cache_hit),
50         .dout(cache_dout),
51         .valid(cache_valid),
52         .dirty(cache_dirty),
53         .tag(cache_tag)
54     );
55
56     localparam
57         S_IDLE = 0,
58         S_PRE_BACK = 1,
59         S_BACK = 2,
60         S_FILL = 3,
61         S_WAIT = 4;
62
63     reg [2:0] state = 0;
64     reg [2:0] next_state = 0;
65     reg [ELEMENT_WORDS_WIDTH-1:0] word_count = 0;
66     reg [ELEMENT_WORDS_WIDTH-1:0] next_word_count = 0;
67     assign cmu_state = state;
68
69     always @ (posedge clk) begin
70         if (rst) begin
71             state <= S_IDLE;
72             word_count <= 2'b00;
73         end
74         else begin
75             state <= next_state;
76             word_count <= next_word_count;
77         end
78     end
79
80     // state ctrl
81     always @ (*) begin
82         if (rst) begin

```

```

83         next_state = S_IDLE;
84         next_word_count = 2'b00;
85     end
86     else begin
87         case (state)
88             S_IDLE: begin
89                 if (en_r || en_w) begin
90                     if (cache_hit)
91                         next_state = S_IDLE;
92                     else if (cache_valid && cache_dirty)
93                         next_state = S_PRE_BACK;
94                     else
95                         next_state = S_FILL;
96                 end
97                 next_word_count = 2'b00;
98             end
99
100             S_PRE_BACK: begin
101                 next_state = S_BACK;
102                 next_word_count = 2'b00;
103             end
104
105             S_BACK: begin
106                 if (mem_ack_i && word_count ==
↪ {ELEMENT_WORDS_WIDTH{1'b1}}) // 2'b11 in default case
107                     next_state = S_FILL;
108                 else
109                     next_state = S_BACK;
110
111                 if (mem_ack_i)
112                     next_word_count = word_count + 2'b01;
113                 else
114                     next_word_count = word_count;
115             end
116
117             S_FILL: begin
118                 if (mem_ack_i && word_count ==
↪ {ELEMENT_WORDS_WIDTH{1'b1}})
119                     next_state = S_WAIT;
120                 else
121                     next_state = S_FILL;
122

```



```

123         if (mem_ack_i)
124             next_word_count = word_count + 2'b01;
125         else
126             next_word_count = word_count;
127         end
128
129         S_WAIT: begin
130             next_state = S_IDLE;
131             next_word_count = 2'b00;
132         end
133     endcase
134 end
135 end
136
137 // cache ctrl
138 always @ (*) begin
139     case(state)
140         S_IDLE, S_WAIT: begin
141             cache_addr = addr_rw;
142             cache_load = en_r;
143             cache_edit = en_w;
144             cache_store = 1'b0;
145             cache_u_b_h_w = u_b_h_w;
146             cache_din = data_w;
147         end
148         S_BACK, S_PRE_BACK: begin
149             cache_addr = {addr_rw[ADDR_BITS-1:BLOCK_WIDTH],
↪ next_word_count, {ELEMENT_WORDS_WIDTH{1'b0}}};
150             cache_load = 1'b0;
151             cache_edit = 1'b0;
152             cache_store = 1'b0;
153             cache_u_b_h_w = 3'b010;
154             cache_din = 32'b0;
155         end
156         S_FILL: begin
157             cache_addr = {addr_rw[ADDR_BITS-1:BLOCK_WIDTH],
↪ word_count, {ELEMENT_WORDS_WIDTH{1'b0}}};
158             cache_load = 1'b0;
159             cache_edit = 1'b0;
160             cache_store = mem_ack_i;
161             cache_u_b_h_w = 3'b010;
162             cache_din = mem_data_i;

```

```

163         end
164     endcase
165 end
166 assign data_r = cache_dout;
167
168 // mem ctrl
169 always @ (*) begin
170     case (next_state)
171         S_IDLE, S_PRE_BACK, S_WAIT: begin
172             mem_cs_o = 1'b0;
173             mem_we_o = 1'b0;
174             mem_addr_o = 32'b0;
175         end
176
177         S_BACK: begin
178             mem_cs_o = 1'b1;
179             mem_we_o = 1'b1;
180             mem_addr_o = {cache_tag,
↵ addr_rw[ADDR_BITS-TAG_BITS-1:BLOCK_WIDTH], next_word_count,
↵ {ELEMENT_WORDS_WIDTH{1'b0}}};
181         end
182
183         S_FILL: begin
184             mem_cs_o = 1'b1;
185             mem_we_o = 1'b0;
186             mem_addr_o = {addr_rw[ADDR_BITS-1:BLOCK_WIDTH],
↵ next_word_count, {ELEMENT_WORDS_WIDTH{1'b0}}};
187         end
188     endcase
189 end
190 assign mem_data_o = cache_dout;
191
192 assign stall = (next_state != S_IDLE);
193
194 endmodule

```

3.2 数据记录及结果分析

Since there are many cases in simulation wave, I only take some examples to show. For the convenience of explanation, all my explanations below are for the positions corresponding to the **yellow vertical lines** in the simulation diagram.

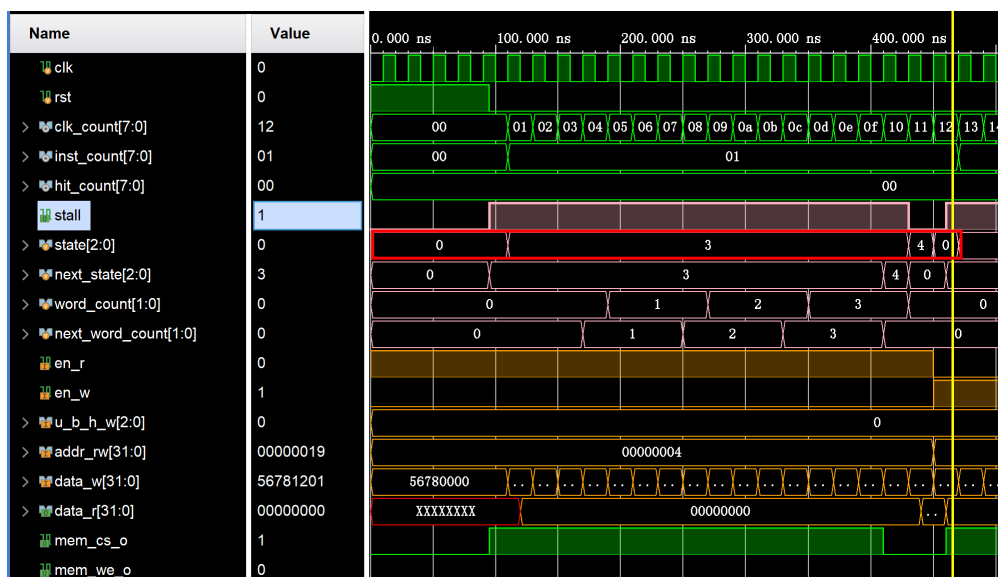
Test bench is as follows:

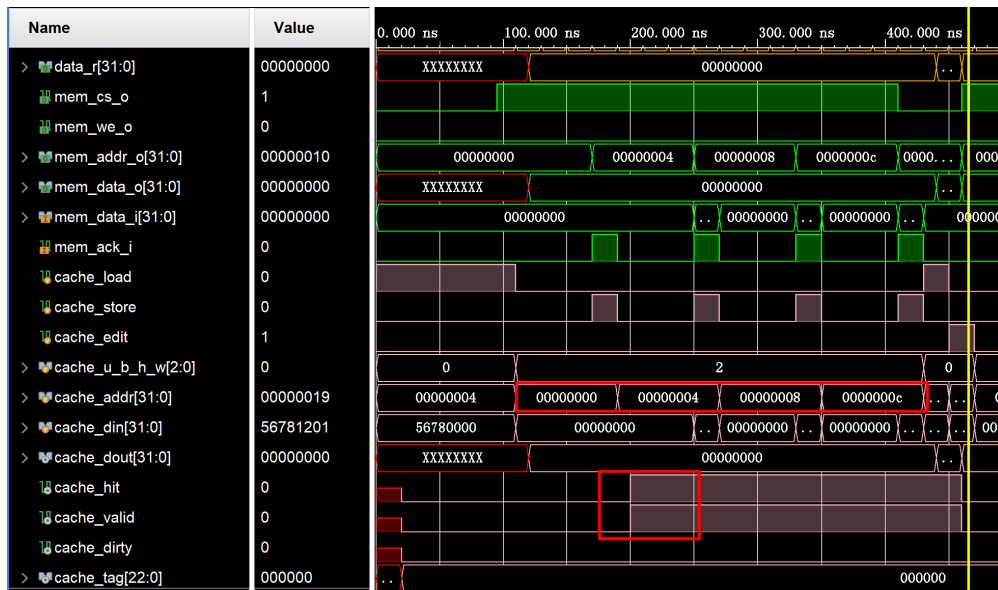
```

initial begin
    data[0] = 40'h0_2_00000004; // read miss          1+17
    data[1] = 40'h0_3_00000019; // write miss         1+17
    data[2] = 40'h1_2_00000008; // read hit            1
    data[3] = 40'h1_3_00000014; // write hit           1
    data[4] = 40'h2_2_00000204; // read miss         1+17
    data[5] = 40'h2_3_00000218; // write miss         1+17
    data[6] = 40'h0_3_00000208; // write hit          1
    data[7] = 40'h4_2_00000414; // read miss + dirty  1+17+17
    data[8] = 40'h1_3_00000404; // write miss + clean 1+17
    data[9] = 40'h0;           // end                total: 128
end

```

- first read miss(clean)





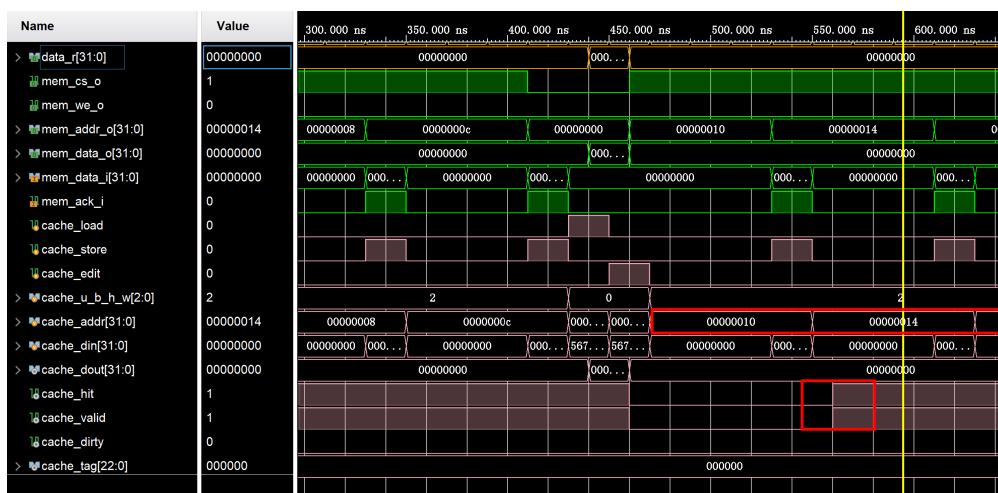
In the corresponding implementation:

```
data[0] = 40'h0_2_00000004; // read miss 1+17
```

Explanation: In the event of a read miss, and assuming there is an available and replaceable position in the cache, the state transitions are as follows: 0 (S_IDLE) → 3 (S_FILL) → 4 (S_WAIT) → 0 (S_IDLE). This corresponds to a normal read miss under the condition of a clean cache, where memory data is written into the cache.

Additionally, during the S_FILL process, because the CPU needs the data at address 0x4, the cache reads the block from memory ranging from 0x0 to 0xC. After reading the data at 0x4, cache_hit and cache_valid become 1.

- first write miss(clean)



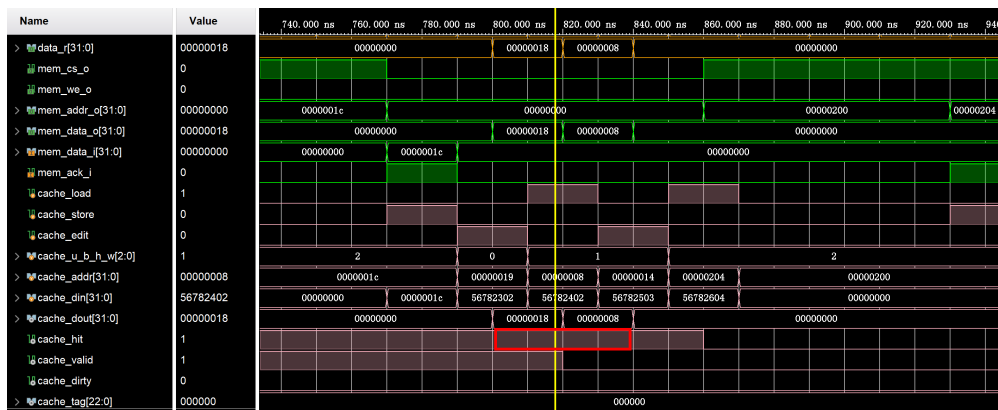
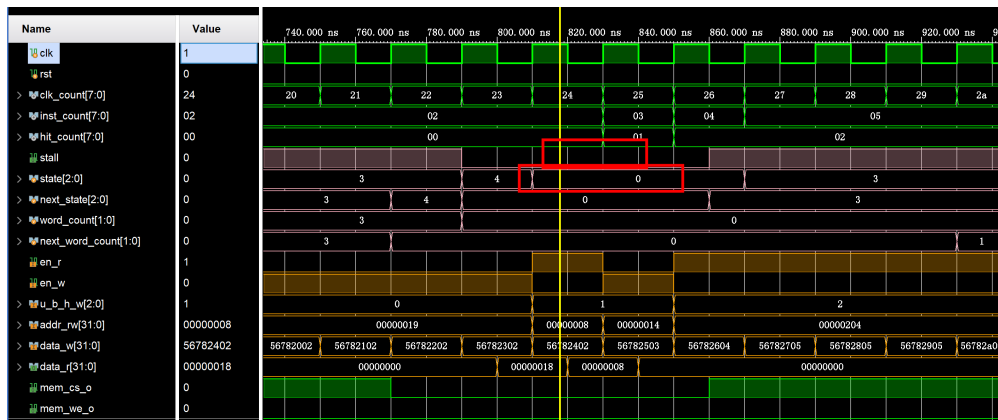


In the corresponding code:

```
data[1] = 40'h0_3_00000019; // write miss 1+17
```

Explanation: The state transitions remain as S_IDLE → S_FILL → S_WAIT → S_IDLE. This represents a write miss under the condition of a clean cache, writing the corresponding data from memory into the cache. Additionally, the write data address is 0x19, so the block from 0x10 to 0x1C is placed into the cache. cache_hit and cache_valid also become 1.

- read hit and write hit

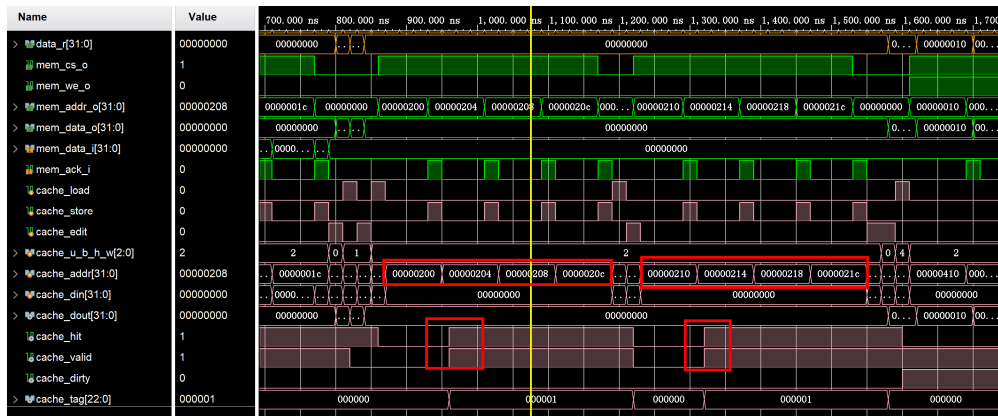
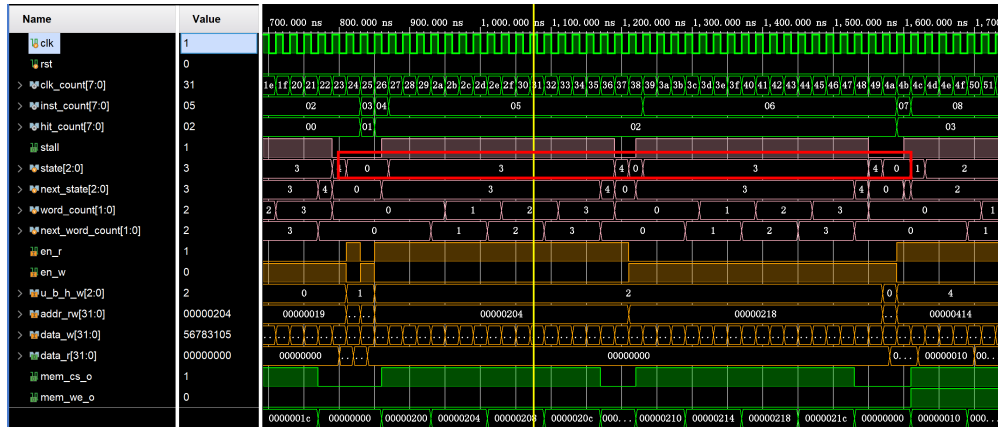


In the corresponding code:

```
data[2] = 40'h1_2_00000008; // read hit 1
data[3] = 40'h1_3_00000014; // write hit 1
```

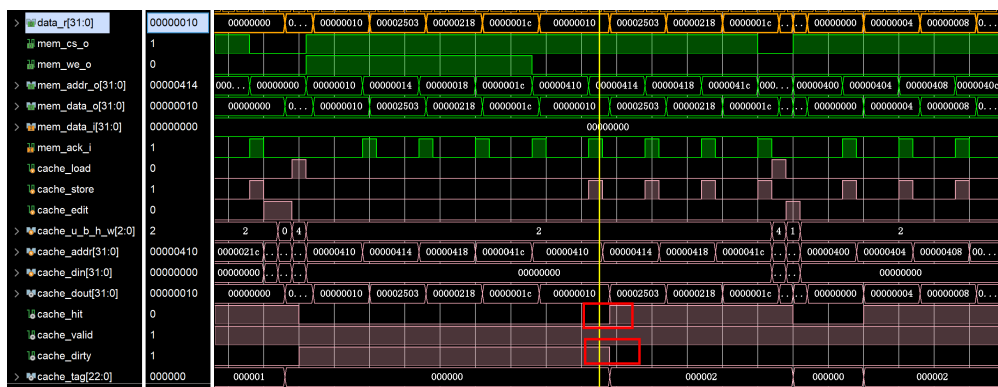
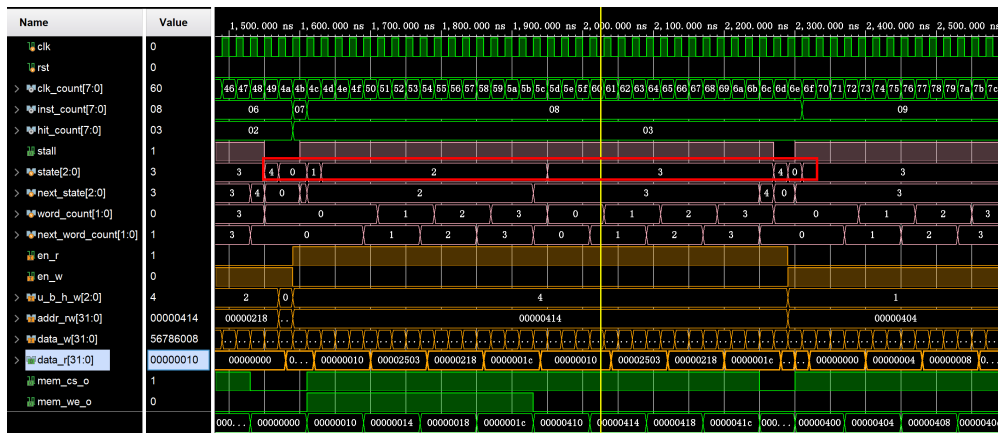
Explanation: When a read hit and a write hit occur, the state machine remains in the S_IDLE state. The CPU does not experience any pause, so the third and fourth instructions execute directly and complete. During this process, stall is 0, and cache_hit is also 1.

- dirty miss replace



```
initial begin
data[4] = 40'h2_2_00000204; // read miss 1+17
data[5] = 40'h2_3_00000218; // write miss 1+17
data[6] = 40'h0_3_00000208; // write hit 1
```

Explanation: In these instructions, the normal replacement process occurs, going through the states S_IDLE → S_FILL → S_WAIT → S_IDLE. The data is loaded into the cache, and cache_hit is set to 1 when the corresponding event occurs. Then, for the write hit, there is no pause, and stall remains 0.

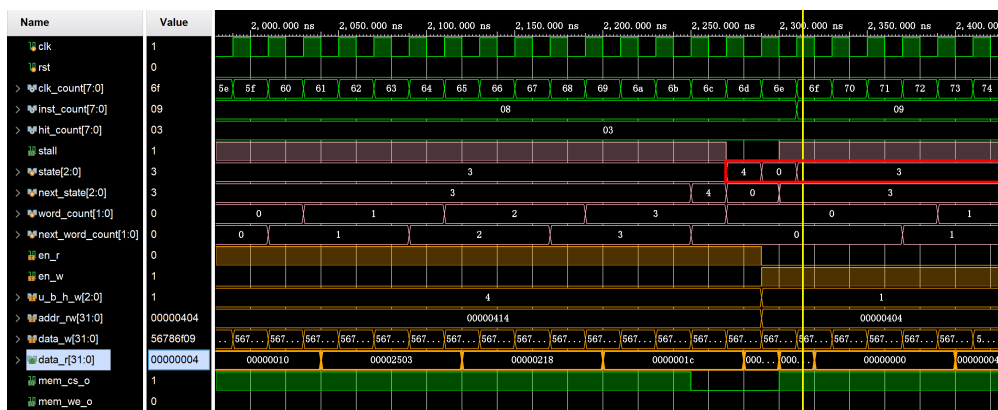


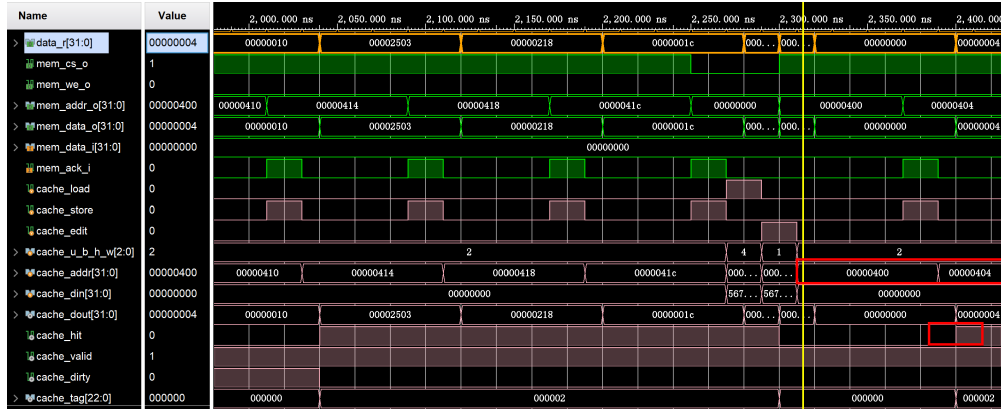
```
data[7] = 40'h4_2_00000414; // read miss + dirty 1+17+17
```

Explanation: To achieve a dirty replacement in a 2-way set-associative cache, it is necessary to fill the cache with the data for the desired operation and modify one of the data blocks to set it as dirty. During the write, this dirty data block is replaced.

The preceding accesses are to create the conditions for this scenario. The state transitions at this point are 0 (S_IDLE) → 1 (S_PRE_BACK) → 2 (S_BACK) → 3 (S_FILL) → 4 (S_WAIT) → 0 (S_IDLE). This sequence corresponds to a dirty replacement in the case of a cache miss, fitting the conditions when cache dirty=1.

- clean miss replace





```
data[8] = 40'h1_3_00000404; // write miss + clean 1+17
```

Explanation: This instruction represents the replacement of a clean block, going through the state transitions $S_IDLE \rightarrow S_FILL \rightarrow S_WAIT \rightarrow S_IDLE$. This process corresponds to a clean replacement in the case of a cache miss.

Now I'll show the figures about the result on sword. Firstly, I'll show data memory and instruction memory.

NO.	Data	Addr.	Comment	NO.	Instruction	Addr.	Comment
0	000080BF	0		16	00000000	40	
1	00000008	4		17	00000000	44	
2	00000010	8		18	00000000	48	
3	00000014	C		19	00000000	4C	
4	FFFF0000	10		20	A3000000	50	
5	0FFF0000	14		21	27000000	54	
6	FF00F0F	18		22	79000000	58	
7	F0F0F0F0	1C		23	15100000	5C	
8	00000000	20		24	00000000	60	
9	00000000	24		25	00000000	64	
10	00000000	28		26	00000000	68	
11	00000000	2C		27	00000000	6C	
12	00000000	30		28	00000000	70	
13	00000000	34		29	00000000	74	
14	00000000	38		30	00000000	78	
15	00000000	3C		31	00000000	7C	

图 1: data memory

NO.	Instruction	Addr.	Label	ASM	Comment
0	00000013	0	__start:	addi x0, x0, 0	
1	01c00083	4		lb x1, 0x01C(x0)	# F0F0F0F0 in 0x1C # miss, read 0x010~0x01C to set 1 line 0
2	01c01103	8		lh x2, 0x01C(x0)	# FFFFFFF0 hit
3	01c02183	C		lw x3, 0x01C(x0)	# F0F0F0F0 hit
4	01c04203	10		lbu x4, 0x01C(x0)	# 000000F0 hit
5	01c05283	14		lhu x5, 0x01C(x0)	# 0000F0F0 hit
6	21002003	18		lw x0, 0x210(x0)	# miss, read 0x210~0x21C to cache set 1 line 1
7	abcde0b7	1C		lw x7, 20(x0)	
8	402200b3	20		lui x1 0xABCDE	
9	71c08093	24		addi x1, x1, 0x71C	# x1 = 0xABCDE71C
10	00100023	28		sb x1, 0x0(x0)	# miss, read 0x000~0x00C to cache set 0 line 0
11	00101223	2C		sh x1, 0x4(x0)	# hit
12	00102423	30		sw x1, 0x8(x0)	# hit

图 2: instruction memory 1

NO.	Instruction	Addr.	Label	ASM	Comment
13	20002303	34		lw x6, 0x200(x0)	# miss, read 0x200~0x20C to cache set 0 line 1
14	40002383	38		lw x7, 0x400(x0)	# miss, write 0x000~0x00C back to ram, then read 0x400~40C to cache set 0 line 0
15	41002403	3C		lw x8, 0x410(x0)	# miss, no write back because of clean, read 0x410~41C to cache set 1 line 0
16	0ed06813	40	loop:	ori x16, x0, 0xED	# end
17	ffdf06f	44		jal x0, loop	

图 3: instruction memory 2

Then I'll show the results.

Zhejiang University Computer Organization Experimental SOC Test Environment (With RISC-V)			
x0: zero 00000000	x01: ra FFFFFFFF	x02: sp 00000000	x03: gp 00000000
x04: tp 00000000	x05: t0 00000000	x06: t1 00000000	x07: t2 00000000
x08: ps0 00000000	x09: s1 00000000	x10: a0 00000000	x11: a1 00000000
x12: a2 00000000	x13: a3 00000000	x14: a4 00000000	x15: a5 00000000
x16: a6 00000000	x17: a7 00000000	x18: s2 00000000	x19: s3 00000000
x20: s4 00000000	x21: s5 00000000	x22: s6 00000000	x23: s7 00000000
x24: s8 00000000	x25: s9 00000000	x26: s10 00000000	x27: s11 00000000
x28: t3 00000000	x29: t4 00000000	x30: t5 00000000	x31: t6 00000000
PC: 1F 00000014	INST-IF 01C05283	rs1Data 00000000	rs2Data 00000000
PC: 1D 00000018	INST-ID 01C04203	rs1Addr 00000000	rs2Addr 00000000
PC: 1E 0000001C	INST-EX 01C02183	CPU-RAM 00000000	PCJump 00000000
PC: 1F 00000018	INST-M 01C01103	B/PCE-S 00000100	D/C-Flag 00000000
PC: 1B 00000004	INST-WB 01C00083	I/ABSel 00010001	PCIFPct 00000000
ALU-Out 00000000	ALU-Out 0000001C	CPUAddr 00000000	ALUCtrl 00000000
ALU-Bit 0000001C	WB-Data FFFFFFFF	CPU-DAI FFFFFFFF	WB-010 00000000
WB-Data 0000001C	WB-Addr 00000001	CPU-DAO 00000000	RegId 00 00000000
CODE-00 00000000			
CODE-01 00000000			
CODE-02 00000000			
CODE-03 01C02183			
CODE-04 01C04203			
CODE-05 00000000			
CODE-06 00000000			
CODE-07 00000000			
CODE-08 00000000			
CODE-09 00000000			
CODE-10 00000000			
CODE-11 00000000			
CODE-12 00000000			
CODE-13 00000000			
CODE-14 00000000			
CODE-15 00000000			
CODE-16 00000000			
CODE-17 00000000			
CODE-18 00000000			
CODE-19 00000000			
CODE-1A 00000000			
CODE-1B 00000000			
CODE-1C 00000000			
CODE-1D 00000000			
CODE-1E 00000000			
CODE-1F 00000000			
CODE-20 00000000			
CODE-21 00000000			
CODE-22 00000000			
CODE-23 00000000			
CODE-24 00000000			
CODE-25 00000000			
CODE-26 00000000			
CODE-27 00000000			

图 4: x1 load correct



图 5: x2 and x3 load correct

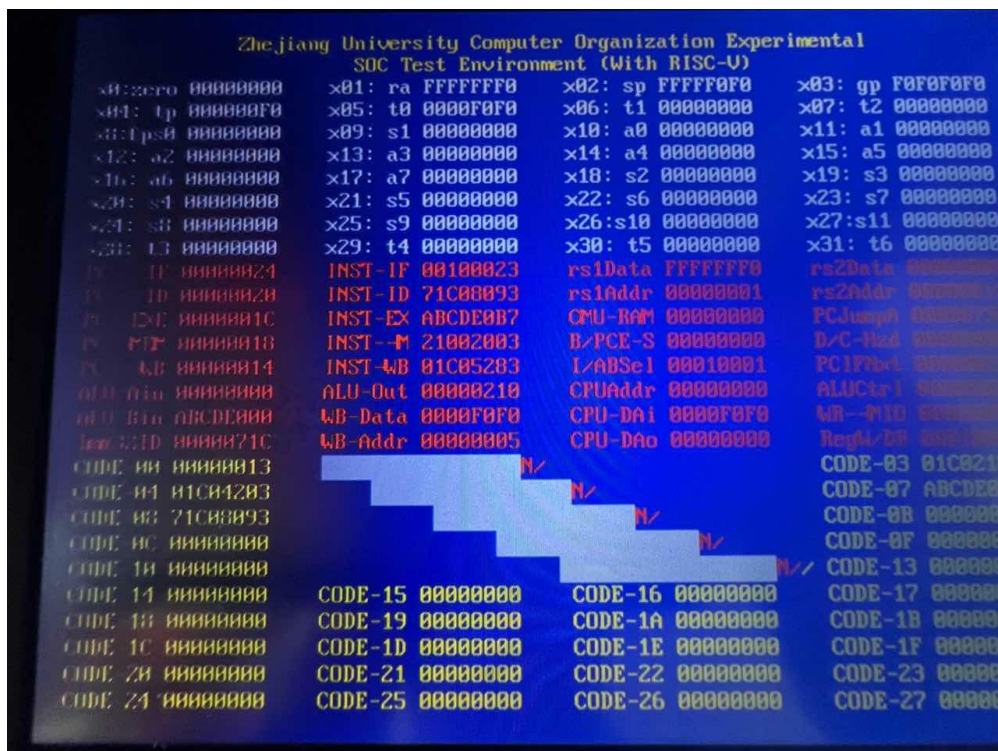


图 6: x4 and x5 load correct

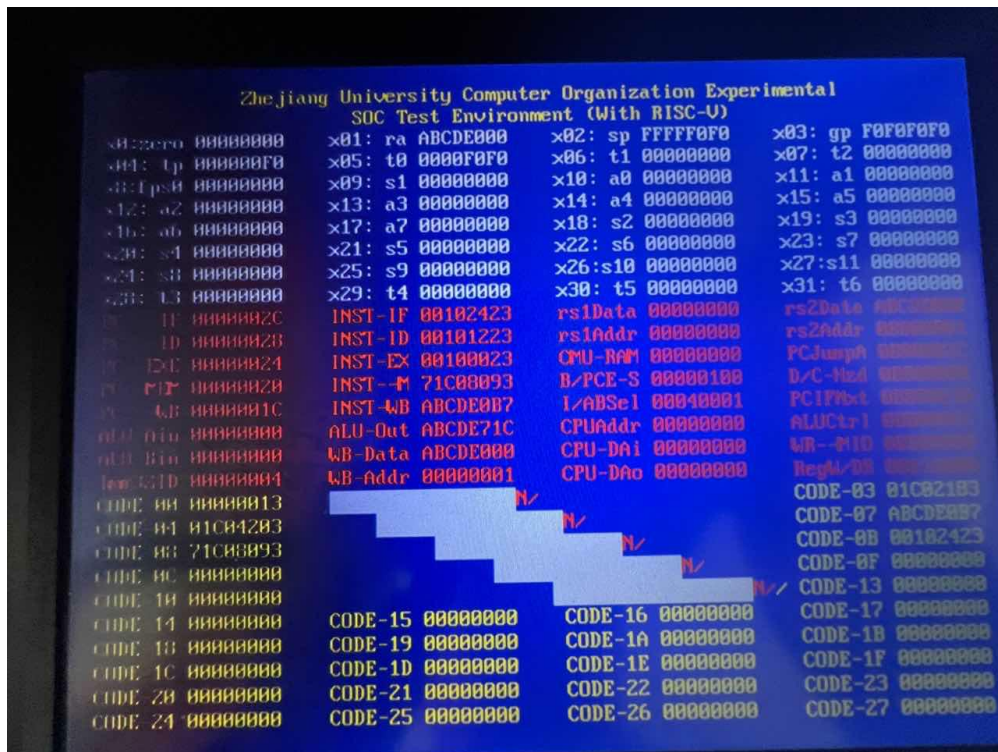


图 7: x1 set to correct value



图 8: x1 add to correct value




图 9: miss and stall

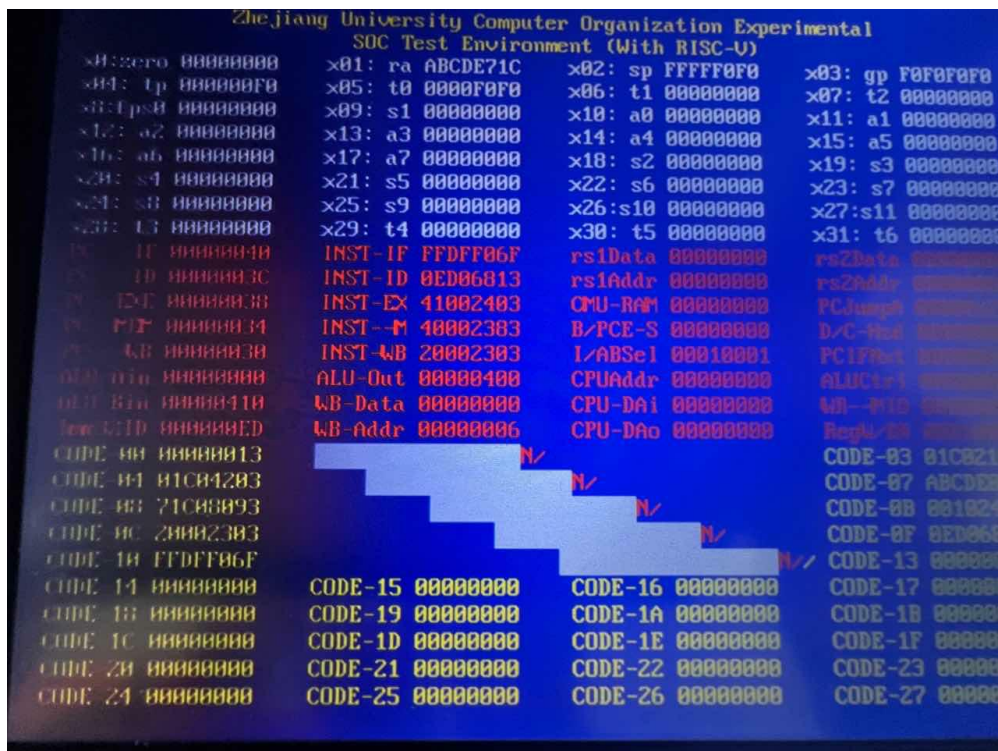


图 10: miss and stall 2

4 讨论与心得

The code need to write for the experiment does not exceed In addition, this experiment requires no more than 10 lines of code written by myself, and the state machine is also very clear. No major difficulties were encountered during the experiment.